

1. We will add a single attribute *ptr* to each node.

For node x , $x.ptr$ leads to a linked list item. The linked list is a circular doubly linked list which contains the labels to all elements in x 's set.

PRINT-SET(x) hence requires constant time, i.e $O(1)$ time to get the list of all elements in x 's set. Since printing each element is $O(1)$, the PRINT-SET(x) operation is linear in the number of elements in the set of x .

The addition of this pointer affects the other operations as follows:

- MAKE-SET(x):

In addition to usual MAKE-SET(x) operations, this implementation will require we create a new linked list, set the label of the linked list item to x and store the pointer of the linked list object in $x.ptr$.

These operations are constant time, i.e $O(1)$ so this does not affect the asymptotic runtime of MAKE-SET(x)

- FIND-SET(x):

In this implementation, FIND-SET(x) does not involve any operations to the linked list hence its asymptotic runtime is not affected.

- UNION(x, y):

Assuming you are merging set y into x , the extra operations in this implementation involve appending the linked list of y to x , which is constant time in a circular doubly linked list. After the linked lists are merged, each node will point to the linked list item which has its label.

2. (a) Let T be a minimum spanning tree. Let (u, v) be an edge in this tree, with weight $w(u, v)$ which is the max weight of all the edges in the MST. Let S be the set of all vertices in T , of which s_1 is the set of vertices whose 'next hop' from u is v . Let s_2 be $S - s_1$. Since a minimum spanning tree cannot have cycles, s_1 and s_2 have to be connected by a single edge, (u, v) , which has to be minimum cost. Hence there is no other spanning tree which can connect the nodes with max edge weight being less than $w(u, v)$. This proves that no bottleneck spanning tree with value less than $w(u, v)$ can exist. Hence a minimum spanning tree is a bottleneck spanning tree.
- (b) To check if a graph G has a bottleneck spanning tree with value b in a linear time algorithm we can use the following algorithm:

1. For each edge E in G
 - 1.1 If edge E has weight greater than b remove it from Graph
2. Let Q be the set of all vertices in G
3. Let S be a queue with a single vertex from Q
4. While S is not empty
 - 4.1 Pop a random vertex V from S
 - 4.2 Remove node V from Q
 - 4.3 For Add all neighbors N of V that exist in Q
 - 4.3.1 Push N into S , if they don't already exist
5. If set Q is empty return true else false

The above algorithm will be linear in number of edges, since each edge is checked once against weight b , then again once when existence of the Spanning Tree is checked.

- (c) We could use the following algorithm on input Graph G :
1. Let E be the edges in graph G .
 2. Let W be the median of weights of $|E|$
 3. Run the Algorithm in 2(b) with G & W
 - 3.1 If a bottleneck spanning tree exists, remove all edges with weight $> W$, and repeat this algorithm with G' .
 - 3.2 If a bottleneck spanning tree does not exist, use MST-Reduce on the forest of edges with weights less than median, and run this algorithm again on G' .

The above will be linear in the number of edges as:

$$O(E + E/2 + E/4 + E/8 + \dots + 1) = O(E)$$

3. (a) In the proposed algorithm, we track usp status for every vertex with an additional array data structure which contains a true or false value about whether a path is unique shortest path.

Initially the usp of each vertex is initialized to true. The usp of source node is assumed as true.

The usp status of a vertex is modified in one of two cases:

- Path with same length as shortest path is found:
We set the usp of vertex to false.
- New Shortest Path:
When a new shortest path is found we set the usp status of vertex to the usp status to the immediate parent of the vertex in the new path.

- (b) The pseudo-code is as follows:

```
function Dijkstra(Graph, source):
  for each vertex v in Graph:
    dist[v] := infinity
    previous[v] := undefined
    usp[v] := true
  dist[source] := 0
  Q := the set of all nodes in Graph
  while Q is not empty:
    u := node in Q with smallest dist[ ]
    remove u from Q
    for each neighbor v of u:
      alt := dist[u] + dist_between(u, v)
      if alt == dist[v]
        usp[v] = false
      if alt < dist[v]
        dist[v] := alt
        previous[v] := u
        usp[v] = usp[u]
  return (previous[ ], dist[ ], usp[ ])
```

- (c) Algorithm run on graph of 5 vertices:

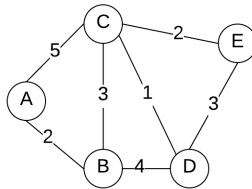


Figure 1: A graph with 5 vertices

Step	Q	B			C			D			E		
		dist	prev	usp	dist	prev	usp	dist	prev	usp	dist	prev	usp
0	A,B,C,D,E	∞	-	T	∞	-	T	∞	-	T	∞	-	T
1	B,C,D,E	2	A	T	5	A	T	∞	-	T	∞	-	T
2	C,D,E	2	A	T	5	A	F	6	B	T	∞	-	T
3	D,E	2	A	T	5	A	F	6	B	F	7	C	F
4	E	2	A	T	5	A	F	6	B	F	7	C	F
5	-	2	A	T	5	A	F	6	B	F	7	C	F

(d) Correctness Proof of *usp*:

From the correctness proof of Dijkstra's Algorithm¹ we use notation $\delta(s, u)$ to define a shortest path between s and u , $u.d$ to define the numeric value of the sum of weights of $\delta(s, u)$ and the conclusion that if u is added to S then $u.d = \delta(s, u)$.

By definition, if any new path to u is found which is equal to $u.d$ then $usp(u) = false$.

Now $usp(u)$ depends on the *usp* status of all the nodes in $\delta(s, u)$. For a vertex y , if $usp(y)$ is false then all v such that $\delta(s, v)$ passes through y , $usp(v) = false$. Also, all items v in path $\delta(y, u)$, $usp(v) = false$.

We use the observation from correctness proof of Dijkstra's Algorithm, that if u is from outside S then the discovered path is minimum distance to u , to prove that if u is added to S , then $usp(parent(u))$ cannot change, i.e since $parent(u).d < u.d$, and $parent(u)$ lies in S .

Hence $usp(u) = usp(parent(u))$ is a valid simplification to the rule where all vertices between u and s have to be checked for being *usp* to figure out $usp(u)$.

(e) The runtime of Dijkstra's Algorithm is not affected by adding the *usp* rules, since creating the *usp* array is at worst $O(v)$, and setting the status of *usp* for every e is $O(1)$; which are both asymptotically slower than $\Theta(E \lg V)$.

Hence time bound of this algorithm is $\Theta(E \lg V)$.

¹Cormen, T. H., Leiserson, C. E., Rivest, R. L., Stein, C. (2009). Introduction to algorithms (3rd ed.). MIT Press.